

# Krish Naresh Patil

Portfolio | LinkedIn | krishpatil4151@gmail.com | (480) 791-9504 | Tempe, AZ

## Summary

Data Engineer with 2+ years building scalable, cloud-native pipelines on GCP and AWS, delivering ETL systems processing 50M+ daily records at Capgemini. Experienced across Airflow, Kafka, Spark, and BigQuery; completing M.S. at ASU, available May 2026.

## Skills

**Programming Languages:** Python, SQL, JavaScript, Java, Bash

**Data Engineering & ETL:** Apache Airflow, Apache Kafka, PySpark, Spark Structured Streaming, dbt, Databricks, Snowflake

**Cloud Platforms:** AWS (EC2, Lambda, Glue, S3, SNS, CloudWatch), GCP (BigQuery, Cloud Storage, Dataflow), Microsoft Fabric

**Data Analytics & Visualization:** Pandas, Scikit-learn, Power BI, Tableau, Looker Studio, SQL Analytics

**Databases:** PostgreSQL, MySQL, MongoDB, Neo4j (Graph DB)

**DevOps & Orchestration:** Docker, Kubernetes, CI/CD, Git

**Tools & Frameworks:** JIRA, Agile/Scrum, FastAPI, REST APIs, Prometheus, Grafana

## Work Experience

**Arizona State University**

Jan 2026 – Present

**Teaching Assistant – Big Data Analytics**

- Designed and delivered lab instruction on Spark, Hadoop, and distributed systems concepts for 100+ graduate students.
- Analyzed 50K+ App Store and Google Play records using Python and SQL to study global platform adoption trends.
- Developed Tableau dashboards visualizing app adoption trends and user behavior, used for coursework and research presentations.

**Flex Technology Group**

Jun 2025 – Aug 2025

**Data Analytics Intern**

- Created 5+ Power BI dashboards integrating SQL Server and Microsoft Fabric tracking revenue, operations, and sales for 500+ clients.
- Streamlined weekly reporting in Microsoft Fabric using SQL and Power BI, reducing manual analysis by 60% (16 to 6.5 hrs/week).
- Migrated legacy SSRS reports to Power BI with enhanced interactivity and drill-through, improving usability and cutting time by 15%.
- Partnered with sales & operations teams to define 20+ KPIs and build data models, consolidating reporting across 3 business units.

**Capgemini**

Jul 2022 - Jul 2024

**Software Engineer**

*Data Engineering — Client Engagement (GCP, BigQuery, Airflow)*

- Architected ETL pipelines with Python, SQL, Airflow, and BigQuery processing 50M+ daily records, reducing latency from 6 to 4 hrs.
- Migrated 50+ legacy SAS workflows to Python on GCP with parallel processing, cutting average job runtime by 55%.
- Automated data ingestion across PostgreSQL, MySQL, MongoDB, and FTP, reducing manual effort by 40% (15 to 6 hrs/week).
- Deployed Kubernetes workflows with monitoring and auto-healing, achieving 99.9% uptime and cutting incident resolution to 45 min.

*Backend Engineering & Cloud Infrastructure (AWS, Databricks)*

- Engineered FTP/SFTP ingestion pipelines for 10+ server feeds, converting raw data into structured, analytics-ready datasets.
- Refactored AWS Glue, Lambda, and S3 workflows to reduce execution time by 30% and improve end-to-end reliability.
- Implemented real-time AWS SNS alerting to surface pipeline failures earlier, reducing operational downtime by 25%.
- Processed 20M+ records on Databricks clusters, optimizing compute usage and improving pipeline throughput by 35%.

## Projects

**Real-Time Fraud Detection Pipeline** — *GitHub*

Feb 2026 – Mar 2026

- Engineered a fraud detection system using Kafka, Spark Structured Streaming, and PostgreSQL, processing 100 TPS with <15s latency.
- Implemented velocity-based fraud detection and Neo4j graph analytics (PageRank, Louvain) to detect fraud networks.
- Orchestrated pipelines with Airflow and dbt, and monitored systems end-to-end using Prometheus and Grafana dashboards.

**TherapyAI** — *Best Use of AI Personas* — *DevHacks 2025*

Oct 2025 - Nov 2025

- Built an AI therapy companion with LLM personas using Gemini API, FastAPI, and pyttsx3 for human-like voice responses.
- Stored user profiles and conversation history as PDFs in AWS S3, enabling personalized therapy continuity across sessions.
- Won Best Use of AI Personas at DevHacks 2025 among 200+ participants for contextually adaptive therapy interactions.

## Education

**M.S. in Software Engineering** Arizona State University, Tempe, AZ

Aug 2024 - May 2026

*GPA: 3.67/4.0 Relevant Coursework: Cloud Computing, Advanced Data Structures & Algorithms, Info Assurance & Security*

**B.E. in Computer Engineering** University of Mumbai, India

Aug 2018 - May 2022

*GPA: 3.46/4.0 Relevant Coursework: Data Structures, Data Warehousing, Machine Learning, Big Data, Distributed Computing*